

Machine Learning with NumPy

Systems Programming
Department of Computer Science and Engineering
Sogang University

Machine Learning Terms

1

Machine learning terms

- **Model:** The process of training an ML model involves providing an ML algorithm (that is, the *learning algorithm*) with training data to learn from. The term *ML model* refers to the model artifact that is created by the training process.
- **Epoch:** An epoch is a single step in training a neural network; in other words when a neural network is trained on every training samples only in one pass we say that one epoch is finished.
- **Learning rate:** The learning rate is how quickly a network abandons old beliefs for new ones.
- **Training dataset:** A training dataset is a dataset of examples used for learning.
- **Test dataset:** A test dataset is a dataset that is independent of the training dataset.


AND/OR Problem

Use “01.AND_OR.ipynb”

```
[ ] # input
x = np.array([[1,1], [1,-1], [-1,-1], [-1,1]])
# output
y = np.array([[1], [0], [0], [0]])

# plot the training data
fig, ax = plt.subplots()
for i in range(y.shape[0]):
    if y[i][0] == 0:
        marker = 'ro'
    else:
        marker = 'bo'
    ax.plot(x[i][0], x[i][1], marker)
ax.axhline(y=0, color='k')
ax.axvline(x=0, color='k')
```

<matplotlib.lines.Line2D at 0x7f85bb925d30>




1

AND Problem

```
[ ] epoch = 5000 # number of training iterations
    learning_rate = 0.1

    # dimension of each layer
    d_in = x.shape[1] # number of features in the input dataset
    d_out = 1 # output layer

    # weight and bias initialization
    wout = np.random.uniform(size=(d_in, 1))
    bout = np.random.uniform(size=(1, d_out))
```



```
[ ] for i in range(epoch):  
    # Forward pass  
    y_pred = sigmoid(x.dot(wout) + bout)  
  
    # Compute and print loss  
    loss = np.square(y_pred - y)  
    if i % 500 == 0:  
        print('Epoch', i, ':', loss.sum())  
  
    # Backpropagation to compute gradients  
    grad_y_pred = (y - y_pred) * derivative_sigmoid(y_pred)  
    grad_wout = x.T.dot(grad_y_pred)  
    grad_bout = np.sum(grad_y_pred, axis=0, keepdims=True)  
  
    # Update weights and biases  
    wout += grad_wout * learning_rate  
    bout += grad_bout * learning_rate
```

```
Epoch 0 : 1.038958564403459  
Epoch 500 : 0.04180152112685249  
Epoch 1000 : 0.0195360067040197  
Epoch 1500 : 0.012512283618333338  
Epoch 2000 : 0.009139607753160103  
Epoch 2500 : 0.007174280895659291  
Epoch 3000 : 0.005892814032537559  
Epoch 3500 : 0.0049934002531985735  
Epoch 4000 : 0.004328425076756785  
Epoch 4500 : 0.003817369959568161
```

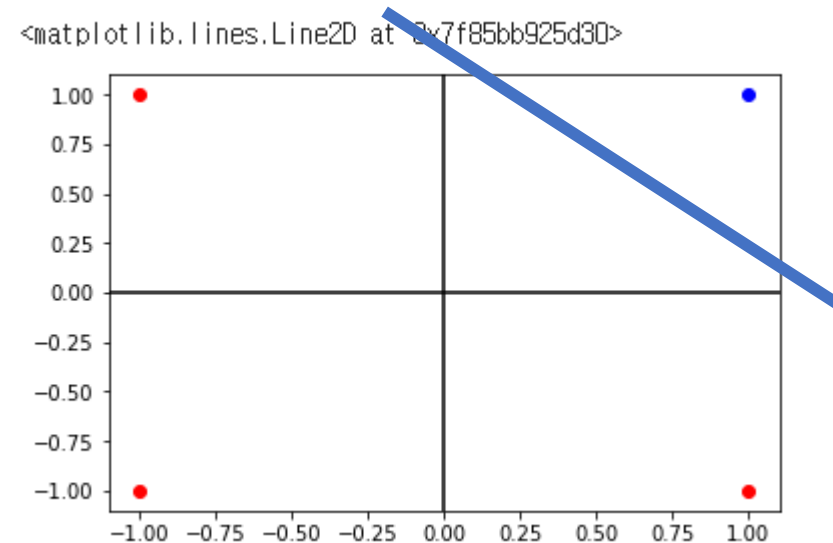
```
[ ] print('Input')
    print(x)
    print('Label')
    print(y)
    print('Output')
    print(y_pred)
    print('Weight')
    print(wout)
    print('Bias')
    print(bout)
```

```
Input
[[ 1  1]
 [ 1 -1]
 [-1 -1]
 [-1  1]]
Label
[[1]
 [0]
 [0]
 [0]]
Output
[[9.66268751e-01]
 [3.37312503e-02]
 [4.25387789e-05]
 [3.37312503e-02]]
Weight
[[3.35512726]
 [3.35512725]]
Bias
[[-3.35512723]]
```



```
Input
[[ 1  1]
 [ 1 -1]
 [-1 -1]
 [-1  1]]
Label
[[1]
 [0]
 [0]
 [0]]
Output
[[9.66268751e-01]
 [3.37312503e-02]
 [4.25387789e-05]
 [3.37312503e-02]]
Weight
[[3.35512726]
 [3.35512725]]
Bias
[[-3.35512723]]
```

$$f(x_1, x_2) = x_1 + x_2 - 1$$




2

OR Problem

```
[ ] # input
x = np.array([[1,1], [1,-1], [-1,-1], [-1,1]])
# output
y = np.array([[1], [1], [0], [1]])

# plot the training data
fig, ax = plt.subplots()
for i in range(y.shape[0]):
    if y[i][0] == 0:
        marker = 'ro'
    else:
        marker = 'bo'
    ax.plot(x[i][0], x[i][1], marker)
ax.axhline(y=0, color='k')
ax.axvline(x=0, color='k')
```

<matplotlib.lines.Line2D at 0x7f85bb89eac8>



2

OR Problem

```
[ ] print('Input')
    print(x)
    print('Label')
    print(y)
    print('Output')
    print(y_pred)
    print('Weight')
    print(wout)
    print('Bias')
    print(bout)
```

```
Input
[[ 1  1]
 [ 1 -1]
 [-1 -1]
 [-1  1]]
Label
[[1]
 [1]
 [0]
 [1]]
Output
[[0.99995794]
 [0.96639196]
 [0.03360804]
 [0.96639196]]
Weight
[[3.35891338]
 [3.35891335]]
Bias
[[3.35891339]]
```


2

AND Problem

```

Input
[[ 1  1]
 [ 1 -1]
 [-1 -1]
 [-1  1]]
Label
[[1]
 [1]
 [0]
 [1]]
Output
[[0.99995794]
 [0.96639196]
 [0.03360804]
 [0.96639196]]
Weight
[[3.35891338]
 [3.35891335]]
Bias
[[3.35891339]]
    
```

<matplotlib.lines.Line2D at 0x7f85bb89eac8>



$$f(x_1, x_2) = x_1 + x_2 + 1$$


XOR Problem

Use “02.XOR.ipynb”

```
[ ] # input
x = np.array([[1,1], [1,-1], [-1,-1], [-1,1]])
# output
y = np.array([[0], [1], [0], [1]])

# plot the training data
fig, ax = plt.subplots()
for i in range(y.shape[0]):
    if y[i][0] == 0:
        marker = 'ro'
    else:
        marker = 'bo'
    ax.plot(x[i][0], x[i][1], marker)
ax.axhline(y=0, color='k')
ax.axvline(x=0, color='k')
```

<matplotlib.lines.Line2D at 0x7fec7e15e320>




2

XOR Problem with Single layer perceptron

```
[ ] epoch = 5000 # number of training iterations
    learning_rate = 0.1

    # dimension of each layer
    d_in = x.shape[1] # number of features in the input dataset
    d_out = 1 # output layer

    # weight and bias initialization
    wout = np.random.uniform(size=(d_in, 1))
    bout = np.random.uniform(size=(1, d_out))
```



2

XOR Problem with Single layer perceptron

```
[ ] for i in range(epoch):  
    # Forward pass  
    y_pred = sigmoid(x.dot(wout) + bout)  
  
    # Compute and print loss  
    loss = np.square(y_pred - y)  
    if i % 500 == 0:  
        print('Epoch', i, ':', loss.sum())  
  
    # Backpropagation to compute gradients  
    grad_y_pred = (y - y_pred) * derivative_sigmoid(y_pred)  
    grad_wout = x.T.dot(grad_y_pred)  
    grad_bout = np.sum(grad_y_pred, axis=0, keepdims=True)  
  
    # Update weights and biases  
    wout += grad_wout * learning_rate  
    bout += grad_bout * learning_rate
```

```
Epoch 0 : 1.1509758455517083  
Epoch 500 : 1.00000000000324432  
Epoch 1000 : 1.0  
Epoch 1500 : 1.0  
Epoch 2000 : 1.0  
Epoch 2500 : 1.0  
Epoch 3000 : 1.0  
Epoch 3500 : 1.0  
Epoch 4000 : 1.0  
Epoch 4500 : 1.0
```


2

XOR Problem with Single layer perceptron

```
[ ] print('Input')
    print(x)
    print('Label')
    print(y)
    print('Output')
    print(y_pred)
    print('Weight')
    print(wout)
    print('Bias')
    print(bout)
```

```
Input
[[ 1  1]
 [ 1 -1]
 [-1 -1]
 [-1  1]]
Label
[[0]
 [1]
 [0]
 [1]]
Output
[[0.5]
 [0.5]
 [0.5]
 [0.5]]
Weight
[[4.06360752e-18]
 [2.32443545e-16]]
Bias
[[-7.17606601e-17]]
```


3

XOR Problem with Multi layer perceptron

```
[ ] epoch = 5000 # number of training iterations
    learning_rate = 0.1

    # dimension of each layer
    d_in = x.shape[1] # number of features in the input dataset
    d_h = 2 # hidden layer
    d_out = 1 # output layer

    # weight and bias initialization
    wh = np.random.uniform(size=(d_in, d_h))
    bh = np.random.uniform(size=(1, d_h))
    wout = np.random.uniform(size=(d_h, d_out))
    bout = np.random.uniform(size=(1, d_out))
```



3

XOR Problem with Multi layer perceptron

```
[ ] for i in range(epoch):  
    # Forward pass  
    h = sigmoid(x.dot(wh) + bh)  
    y_pred = sigmoid(h.dot(wout) + bout)  
  
    # Compute and print loss  
    loss = np.square(y_pred - y)  
    if i % 500 == 0:  
        print('Epoch', i, ':', loss.sum())  
  
    # Backpropagation to compute gradients  
    grad_y_pred = (y - y_pred) * derivative_sigmoid(y_pred)  
    grad_wout = h.T.dot(grad_y_pred)  
    grad_bout = np.sum(grad_y_pred, axis=0, keepdims=True)  
    grad_h = grad_y_pred.dot(wout.T) * derivative_sigmoid(h)  
    grad_wh = x.T.dot(grad_h)  
    grad_bh = np.sum(grad_h, axis=0, keepdims=True)  
  
    # Update weights and biases  
    wout += grad_wout * learning_rate  
    bout += grad_bout * learning_rate  
    wh += grad_wh * learning_rate  
    bh += grad_bh * learning_rate
```

```
Epoch 0 : 1.264783292711578  
Epoch 500 : 1.000410301966061  
Epoch 1000 : 0.9997919590083908  
Epoch 1500 : 0.9989677933762147  
Epoch 2000 : 0.995787723453303  
Epoch 2500 : 0.9742161349844749  
Epoch 3000 : 0.8669798788416558  
Epoch 3500 : 0.7475512915835518  
Epoch 4000 : 0.503842729214347  
Epoch 4500 : 0.1408154431024505
```

3

XOR Problem with Multi layer perceptron

```
[ ] print('Input')
    print(x)
    print('Label')
    print(y)
    print('Output')
    print(y_pred)
    print('Weight @ Hidden layer')
    print(wh)
    print('Bias @ Hidden layer')
    print(bh)
    print('Weight @ Output layer')
    print(wout)
    print('Bias @ Output layer')
    print(bout)
```

```
Input
[[ 1  1]
 [ 1 -1]
 [-1 -1]
 [-1  1]]
Label
[[0]
 [1]
 [0]
 [1]]
Output
[[0.12470714]
 [0.86727314]
 [0.12523307]
 [0.86244534]]
Weight @ Hidden layer
[[-2.40124599 -3.25673262]
 [ 2.40817    3.29521572]]
Bias @ Hidden layer
[[-2.22600371  3.45598513]]
Weight @ Output layer
[[ 4.7150911 ]
 [-4.62109687]]
Bias @ Output layer
[[2.07336382]]
```